

1nt19is156

Program 3 : Map reduce program for word count

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class WordCount {
    public static class TokenizerMapper
        extends Mapper{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(Object key, Text value, Context context ) throws IOException, InterruptedException
        { StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        } }
    public static class IntSumReducer extends Reducer
    {
        private IntWritable result = new IntWritable();
        public void reduce(Text key, Iterable values,
            Context context ) throws IOException, InterruptedException {
            int sum = 0;
```

```

for (IntWritable val : values)
{ sum += val.get();
} result.set(sum);
context.write(key, result); } }

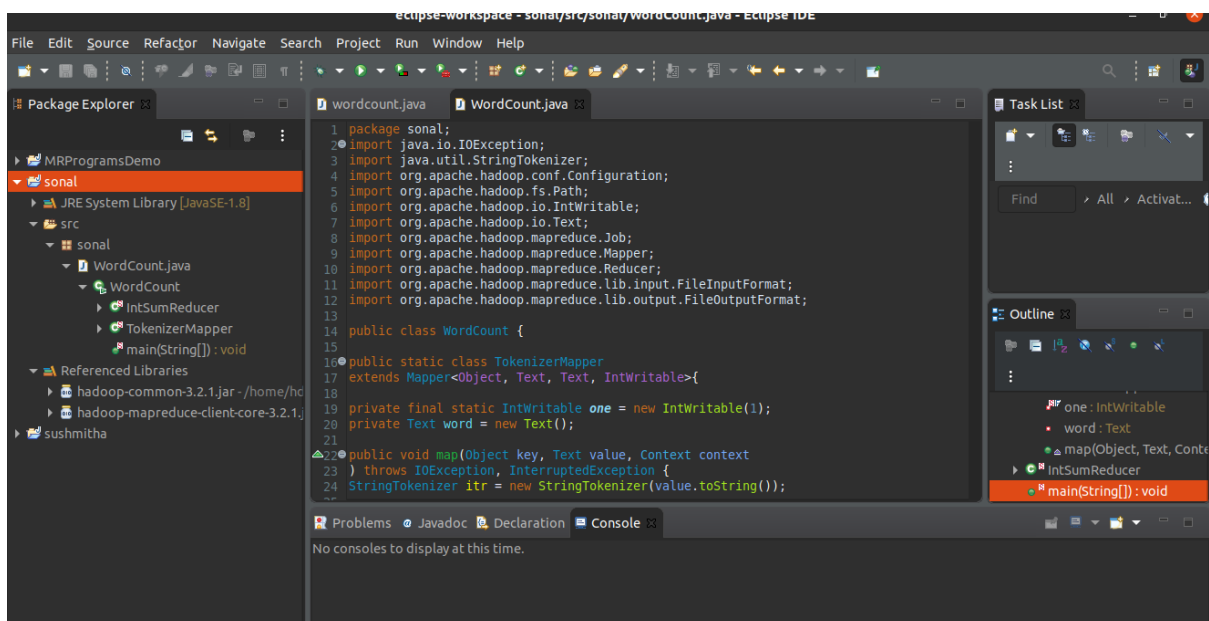
public static void main(String[] args)
throws Exception { Configuration conf = new Configuration();

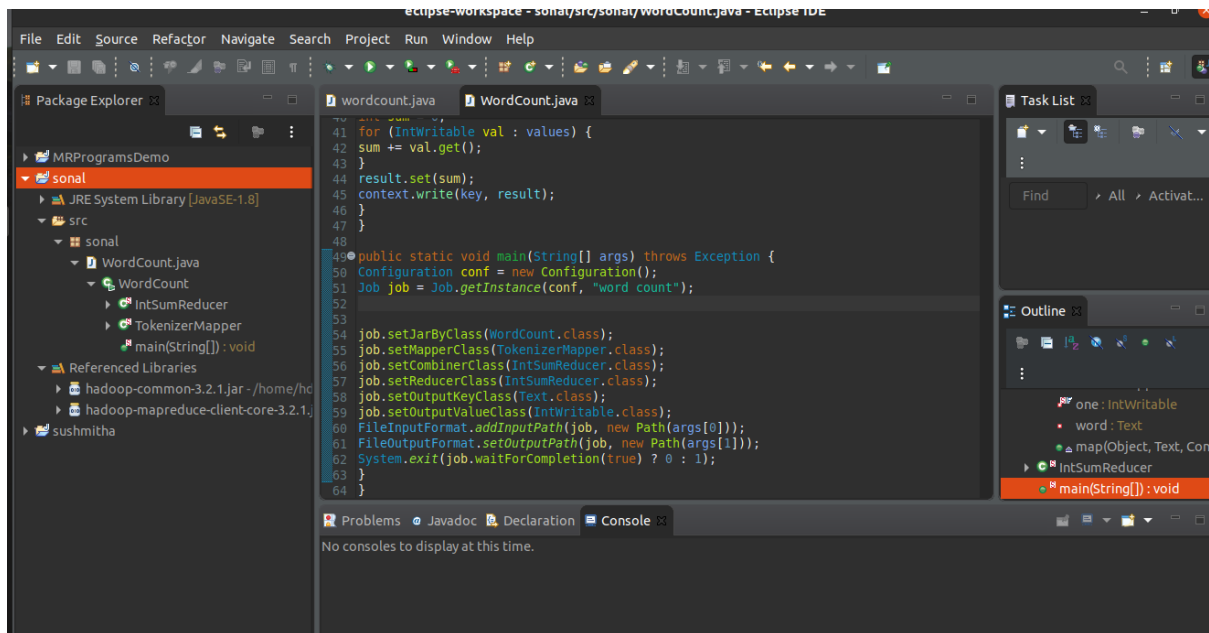
Job job = Job.getInstance(conf, "word count");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);

FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));

System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```





```
hadoop@admin1-HP-280-G4-MT-Business-PC:~$ cd $HADOOP_HOME
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1$ cd sbin
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ jps
4017 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar
5796 NodeManager
5063 NameNode
5641 ResourceManager
8475 Jps
5211 DataNode
5453 SecondaryNameNode
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 5063. Stop it first.
Starting datanodes
localhost: datanode is running as process 5211. Stop it first.
Starting secondary namenodes [admin1-HP-280-G4-MT-Business-PC]
admin1-HP-280-G4-MT-Business-PC: secondarynamenode is running as process 5453. Stop it first.
Starting resourcemanager
resourcemanager is running as process 5641. Stop it first.
Starting nodemanagers
localhost: nodemanager is running as process 5796. Stop it first.
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -mkdir -p ~/input
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -appendToFile ~/input/test.txt
hy there how are you I am fine 2022-06-17 10:38:44,216 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, r
emoteHostTrusted = false
```

```
at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hadoop jar /home/hadoop/Desktop/sonal.jar ~/input ~/out
2022-06-17 10:46:34,549 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-17 10:46:34,781 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and ex
ecute your application with ToolRunner to remedy this.
2022-06-17 10:46:34,803 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_165544
0953471_0002
2022-06-17 10:46:34,884 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-17 10:46:34,982 INFO input.FileInputFormat: Total input files to process : 1
2022-06-17 10:46:35,095 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-17 10:46:35,553 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-17 10:46:35,569 INFO mapreduce.JobSubmitter: number of splits:1
2022-06-17 10:46:35,686 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-17 10:46:36,111 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1655440953471_0002
2022-06-17 10:46:36,111 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-06-17 10:46:36,246 INFO conf.Configuration: resource-types.xml not found
2022-06-17 10:46:36,246 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-06-17 10:46:36,420 INFO impl.YarnClientImpl: Submitted application application_1655440953471_0002
2022-06-17 10:46:36,452 INFO mapreduce.Job: The url to track the job: http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1655440953471
_0002/
2022-06-17 10:46:36,452 INFO mapreduce.Job: Running job: job_1655440953471_0002
2022-06-17 10:46:42,549 INFO mapreduce.Job: Job job_1655440953471_0002 running in uber mode : false
2022-06-17 10:46:42,551 INFO mapreduce.Job: map 0% reduce 0%
2022-06-17 10:46:45,602 INFO mapreduce.Job: map 100% reduce 0%
2022-06-17 10:46:49,628 INFO mapreduce.Job: map 100% reduce 100%
2022-06-17 10:46:49,650 INFO mapreduce.Job: Job job_1655440953471_0002 completed successfully
2022-06-17 10:46:49,730 INFO mapreduce.Job: Counters: 54
File System Counters
FILE: Number of bytes read=98
FILE: Number of bytes written=451079
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=168
HDFS: Number of bytes written=56
HDFS: Number of read operations=8
HDFS: Number of large read operations=0
```

```

File System Counters
  FILE: Number of bytes read=98
  FILE: Number of bytes written=451079
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=168
  HDFS: Number of bytes written=56
  HDFS: Number of read operations=8
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=1385
  Total time spent by all reduces in occupied slots (ms)=1497
  Total time spent by all map tasks (ms)=1385
  Total time spent by all reduce tasks (ms)=1497
  Total vcore-milliseconds taken by all map tasks=1385
  Total vcore-milliseconds taken by all reduce tasks=1497
  Total megabyte-milliseconds taken by all map tasks=1418240
  Total megabyte-milliseconds taken by all reduce tasks=1532928
Map-Reduce Framework
  Map input records=1
  Map output records=13
  Map output bytes=108
  Map output materialized bytes=98
  Input split bytes=112
  Combine input records=13
  Combine output records=9
  Reduce input groups=9
  Reduce shuffle bytes=98
  Reduce input records=9
  Reduce output records=9
  Spilled Records=18
  Shuffled Maps =1
  Failed Shuffles=0

```

```

Map-Reduce Framework
  Map input records=1
  Map output records=13
  Map output bytes=108
  Map output materialized bytes=98
  Input split bytes=112
  Combine input records=13
  Combine output records=9
  Reduce input groups=9
  Reduce shuffle bytes=98
  Reduce input records=9
  Reduce output records=9
  Spilled Records=18
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=64
  CPU time spent (ms)=710
  Physical memory (bytes) snapshot=485638144
  Virtual memory (bytes) snapshot=5080743936
  Total committed heap usage (bytes)=517996544
  Peak Map Physical memory (bytes)=294223872
  Peak Map Virtual memory (bytes)=2535649280
  Peak Reduce Physical memory (bytes)=191414272
  Peak Reduce Virtual memory (bytes)=2545094656
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=56
File Output Format Counters
  Bytes Written=56

```

output

```

      Bytes Read=56
File Output Format Counters
  Bytes Written=56
hadoop@admin1-HP-220-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -cat ~/out/part*
2022-06-17 10:47:44,841 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
am      1
are     2
fine    1
hello   1
how     2
hy      1
i       1
there   2
you     2

```

Program 4 : Map reduce program

```
import java.io.IOException;
import java.util.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class TransactionCount {
//MAPPER CODE

public static class Map extends MapReduceBase implements
Mapper<LongWritable, Text, Text, IntWritable> {
private final static IntWritable one = new IntWritable(1);
//private Text word = new Text();

public void map(LongWritable key, Text value, OutputCollector<Text,
IntWritable> output, Reporter reporter) throws IOException {
String myString = value.toString();
String[] userCount = myString.split(",");
output.collect(new Text(userCount[3]), one);

}
}

//REDUCER CODE
public static class Reduce extends MapReduceBase implements
Reducer<Text, IntWritable, Text, IntWritable> {
public void reduce(Text key, Iterator<IntWritable> values,
OutputCollector<Text, IntWritable> output, Reporter reporter) throws
IOException { //{little: {1,1}}
int finaluserCount = 0 ;
Text mykey = key ;
while(values.hasNext()) {
IntWritable value = values.next();
finaluserCount += value.get();
}
output.collect(mykey, new IntWritable(finaluserCount));
}
}

//DRIVER CODE
public static void main(String[] args) throws Exception {
JobConf conf = new JobConf(TransactionCount.class);
conf.setJobName("wordcount");
conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(IntWritable.class);
conf.setMapperClass(Map.class);
```

```

conf.setCombinerClass(Reduce.class);
conf.setReducerClass(Reduce.class);
conf.setInputFormat(TextInputFormat.class);
conf.setOutputFormat(TextOutputFormat.class);
FileInputFormat.setInputPaths(conf, new Path(args[0]));
FileOutputFormat.setOutputPath(conf, new Path(args[1]));
JobClient.runJob(conf);
}

}

```

```

hadoop@admin1-HP-280-G4-MT-Business-PC: ~/hadoop-3.2.1/sbin
hadoop@admin1-HP-280-G4-MT-Business-PC:~$ cd $HADOOP_HOME/sbin
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
localhost: namenode is running as process 7044. Stop it first.
Starting datanodes
localhost: datanode is running as process 7199. Stop it first.
Starting secondary namenodes [admin1-HP-280-G4-MT-Business-PC]
admin1-HP-280-G4-MT-Business-PC: secondarynamenode is running as process 7438. Stop it first.
Starting resourcemanager
resourcemanager is running as process 7626. Stop it first.
Starting nodemanagers
localhost: nodemanager is running as process 7776. Stop it first.
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ jps
7776 NodeManager
7044 NameNode
18217 Jps
16634 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar
7626 ResourceManager
7438 SecondaryNameNode
7199 DataNode

```

```

hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -mkdir -p /input156
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -copyFromLocal /home/hadoop/Documents/sonal.csv /input156
2022-06-20 10:57:54,570 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hadoop jar /home/hadoop/Desktop/sonal156.jar /input156 /output156
2022-06-20 10:58:33,497 INFO client.RMPProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-20 10:58:33,619 INFO client.RMPProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-06-20 10:58:33,748 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2022-06-20 10:58:33,771 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1655698442856_0006
2022-06-20 10:58:34,129 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-20 10:58:34,258 INFO mapred.FileInputFormat: Total input files to process : 1
2022-06-20 10:58:34,313 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-20 10:58:34,363 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-20 10:58:34,431 INFO mapreduce.JobSubmitter: number of splits:2
2022-06-20 10:58:34,562 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-06-20 10:58:34,570 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1655698442856_0006
2022-06-20 10:58:34,570 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-06-20 10:58:34,688 INFO conf.Configuration: resource-types.xml not found
2022-06-20 10:58:34,688 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-06-20 10:58:34,730 INFO impl.YarnClientImpl: Submitted application application_1655698442856_0006
2022-06-20 10:58:34,755 INFO mapreduce.Job: The url to track the job: http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1655698442856_0006/

```

Output

```

File Output Format Counters
  Bytes Written=33
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ hdfs dfs -cat /output156/part*
2022-06-20 10:58:59,751 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
10000 1
20000 3
30000 1
amount 1

```